

```
!pip install pymongo
Collecting pymongo
  Using cached pymongo-4.8.0-cp312-cp312-win_amd64.whl.metadata (22
kB)
Collecting dnspython<3.0.0,>=1.16.0 (from pymongo)
  Using cached dnspython-2.6.1-py3-none-any.whl.metadata (5.8 kB)
Using cached pymongo-4.8.0-cp312-cp312-win_amd64.whl (680 kB)
Using cached dnspython-2.6.1-py3-none-any.whl (307 kB)
Installing collected packages: dnspython, pymongo
Successfully installed dnspython-2.6.1 pymongo-4.8.0

[notice] A new release of pip is available: 24.0 -> 24.2
[notice] To update, run: python.exe -m pip install --upgrade pip

import pymongo
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Assignment 4 - MongoDB & PyMongo

Dataset Description:

- ID: Rows ID
- YEAR: 2015
- MONTH: 1-12
- DAY: 1-31
- DAY_OF_WEEK: 1 (Monday) - 7 (Sunday)
- AIRLINE: Airline CODE
- FLIGHT_NUMBER: Flight Number
- TAIL_NUMBER: Flight's tail number
- ORIGIN_AIRPORT: Origin IATA airport code
- DESTINATION_AIRPORT: Destination IATA airport code
- SCHEDULED_DEPARTURE: Actual departure time (local, hhmm)
- DEPARTURE_TIME: Scheduled departure time (local, hhmm)
- DEPARTURE_DELAY: Departure delay, in minutes
- TAXI_OUT: Taxi out time in minutes
- WHEELS_OFF:
- SCHEDULED_TIME: Scheduled arrival time (local, hhmm)
- ELAPSED_TIME: in Minutes
- AIR_TIME: in Minutes
- DISTANCE: in Miles

- WHEELS_ON:
- TAXI_IN: Taxi in time, in minutes
- SCHEDULED_ARRIVAL: Scheduled arrival time (local, hhmm)
- ARRIVAL_TIME: Actual arrival time (local, hhmm)
- ARRIVAL_DELAY: Arrival delay, in minutes
- DIVERTED: 1 = yes, 0 = no
- CANCELLED: 1 = yes, 0 = no

1. Create collections "flights" inside database "airline_delayDB"

```
# connecting to MongoDB
client = pymongo.MongoClient('localhost:27017')
print(client)

databaseName = 'airline_delayDB'
collectionName = 'flights'

# create database if not exists
# create collection if it doesnt exist

db = client[databaseName]
print(db)
collection = db[collectionName]
print(collection)

MongoClient(host=['localhost:27017'], document_class=dict,
tz_aware=False, connect=True)
Database(MongoClient(host=['localhost:27017'], document_class=dict,
tz_aware=False, connect=True), 'airline_delayDB')
Collection(Database(MongoClient(host=['localhost:27017'],
document_class=dict, tz_aware=False, connect=True),
'airline_delayDB'), 'flights')
```

2. How would you insert this entire dataset into a MongoDB collection named flights? Describe the structure of each document.

```
fileLocation = r'Flights_Delay.csv'
airline_df = pd.read_csv(fileLocation)
print(airline_df)

records = airline_df.to_dict(orient='records')

if records:
    collection.insert_many(records)
```

	ID	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	\
0	0	2015	3	4	3	EV	5170	
1	1	2015	2	2	1	MQ	3584	

2	2	2015	1	27	2	B6	716
3	3	2015	1	28	3	EV	4289
4	4	2015	2	5	4	EV	5584
...
55995	55995	2015	2	4	3	B6	1567
55996	55996	2015	1	17	6	AA	1113
55997	55997	2015	1	17	6	US	661
55998	55998	2015	1	13	2	DL	1318
55999	55999	2015	2	25	3	US	499

TAIL_NUMBER ORIGIN_AIRPORT DESTINATION_AIRPORT ...							
ARRIVAL_TIME \							
0	N842AS		CVG		XNA	...	
1103.0							
1	N646MQ		DFW		SPS	...	
1402.0							
2	N309JB		JAX		DCA	...	
1655.0							
3	N14162		COS		IAH	...	
1742.0							
4	N851AS		ATL		AVL	...	
1352.0							
...
.							
55995	N508JB		HPN		PBI	...	
1338.0							
55996	N4YBAA		PIT		DFW	...	
1055.0							
55997	N534UW		LAX		PHL	...	
1424.0							
55998	N348NB		ATL		CLT	...	
1116.0							
55999	NaN		MIA		CLT	...	
NaN							

ARRIVAL_DELAY DIVERTED CANCELLED CANCELLATION_REASON \					
0	33.0	0	0		NaN
1	32.0	0	0		NaN
2	96.0	0	0		NaN
3	-19.0	0	0		NaN
4	9.0	0	0		NaN
...
55995	174.0	0	0		NaN
55996	-30.0	0	0		NaN
55997	-3.0	0	0		NaN
55998	-3.0	0	0		NaN
55999	NaN	0	1		B

AIR_SYSTEM_DELAY SECURITY_DELAY AIRLINE_DELAY			
LATE_AIRCRAFT_DELAY \			

0	14.0	0.0	19.0
0.0			
1	0.0	0.0	32.0
0.0			
2	6.0	0.0	90.0
0.0			
3	NaN	NaN	NaN
NaN			
4	NaN	NaN	NaN
NaN			
...
...			
55995	10.0	0.0	164.0
0.0			
55996	NaN	NaN	NaN
NaN			
55997	NaN	NaN	NaN
NaN			
55998	NaN	NaN	NaN
NaN			
55999	NaN	NaN	NaN
NaN			

	WEATHER_DELAY
0	0.0
1	0.0
2	0.0
3	NaN
4	NaN
...	...
55995	0.0
55996	NaN
55997	NaN
55998	NaN
55999	NaN

[56000 rows x 32 columns]

3. Write a MongoDB command to insert a single flight record from the dataset.

```
oneRecord = airline_df.head(1)
oneRecord = oneRecord.to_dict(orient='records')[0]

if oneRecord:
    collection.insert_one(oneRecord)
```

Queries

4. Write a MongoDB query to find all flights that were delayed by more than 60 minutes.

```
# DEPARTURE_DELAY >= 60
list(
  collection.find(
    {
      'DEPARTURE_DELAY': {'$gte': 60}
    }, {
      '_id':0, 'AIRLINE':1, 'FLIGHT_NUMBER':1, 'TAIL_NUMBER':1
    }
  )
  .limit(10)
)

[{'AIRLINE': 'B6', 'FLIGHT_NUMBER': 716, 'TAIL_NUMBER': 'N309JB'},
 {'AIRLINE': '00', 'FLIGHT_NUMBER': 6196, 'TAIL_NUMBER': 'N751SK'},
 {'AIRLINE': 'US', 'FLIGHT_NUMBER': 1756, 'TAIL_NUMBER': 'N823AW'},
 {'AIRLINE': '00', 'FLIGHT_NUMBER': 2699, 'TAIL_NUMBER': 'N897SK'},
 {'AIRLINE': 'F9', 'FLIGHT_NUMBER': 661, 'TAIL_NUMBER': 'N922FR'},
 {'AIRLINE': 'US', 'FLIGHT_NUMBER': 686, 'TAIL_NUMBER': 'N570UW'},
 {'AIRLINE': '00', 'FLIGHT_NUMBER': 4544, 'TAIL_NUMBER': 'N825SK'},
 {'AIRLINE': 'WN', 'FLIGHT_NUMBER': 1165, 'TAIL_NUMBER': 'N742SW'},
 {'AIRLINE': 'EV', 'FLIGHT_NUMBER': 3936, 'TAIL_NUMBER': 'N14180'},
 {'AIRLINE': 'DL', 'FLIGHT_NUMBER': 1088, 'TAIL_NUMBER': 'N688DL'}]
```

5. How would you query all flights that were cancelled (CANCELLED flag set to 1) and return only the AIRLINE, ORIGIN_AIRPORT, and CANCELLATION_REASON fields?

```
list(
  collection.find(
    {
      'CANCELLED': {'$eq': 1}
    }, {
      '_id':0, 'AIRLINE':1, 'ORIGIN_AIRPORT':1,
      'CANCELLATION_REASON':1
    }
  )
  .limit(10)
)

[{'AIRLINE': 'EV', 'ORIGIN_AIRPORT': 'MLI', 'CANCELLATION_REASON': 'C'},
 {'AIRLINE': 'WN', 'ORIGIN_AIRPORT': 'BWI', 'CANCELLATION_REASON': 'B'},
 {'AIRLINE': 'DL', 'ORIGIN_AIRPORT': 'SFO', 'CANCELLATION_REASON': 'C'}]
```

```
'B'},
{'AIRLINE': 'AA', 'ORIGIN_AIRPORT': 'DFW', 'CANCELLATION_REASON':
'B'},
{'AIRLINE': 'MQ', 'ORIGIN_AIRPORT': 'LGA', 'CANCELLATION_REASON':
'B'},
{'AIRLINE': 'AA', 'ORIGIN_AIRPORT': 'BDL', 'CANCELLATION_REASON':
'B'},
{'AIRLINE': 'WN', 'ORIGIN_AIRPORT': 'MKE', 'CANCELLATION_REASON':
'B'},
{'AIRLINE': 'US', 'ORIGIN_AIRPORT': 'DCA', 'CANCELLATION_REASON':
'B'},
{'AIRLINE': 'WN', 'ORIGIN_AIRPORT': 'FLL', 'CANCELLATION_REASON':
'B'},
{'AIRLINE': 'EV', 'ORIGIN_AIRPORT': 'ORF', 'CANCELLATION_REASON':
'B'}]
```

6. Using MongoDB's aggregation framework, how would you calculate the average arrival delay (ARRIVAL_DELAY) for each airline? [Create a suitable plot using matplotlib/seaborn]

```
# list(collection.distinct('ARRIVAL_DELAY'))
queryResult = list(
    collection.aggregate([
        {'$match': {'ARRIVAL_DELAY': {'$exists': True, '$ne':
float('nan'), '$gte':0}}},
        {'$group': {'_id':'$AIRLINE', 'avgArrivalDelay': {'$avg':
'$ARRIVAL_DELAY'}}},
        {'$project': {
            '_id':0,
            'Airline': '$_id',
            'avgArrivalDelay': {'$round': ['$avgArrivalDelay', 3]}
        }}
    ])
)

for item in queryResult:
    print(item)

df = pd.DataFrame(queryResult)
df.columns = ['Airline', 'AverageArrivalDelay']


plt.figure(figsize=(25,8))
bar_plot = sns.barplot(data=df, x='AverageArrivalDelay', y='Airline')

plt.xlabel('Average Arrival Delay (in minutes)')
plt.ylabel('Airline')
plt.title('Average Arrival Delay for each Airline')
```

```
bar_plot.bar_label(bar_plot.containers[0])
```

```
plt.show()
```

```
{'Airline': 'AS', 'avgArrivalDelay': 23.119}  
{'Airline': 'UA', 'avgArrivalDelay': 32.457}  
{'Airline': 'AA', 'avgArrivalDelay': 34.603}  
{'Airline': 'WN', 'avgArrivalDelay': 26.03}  
{'Airline': 'US', 'avgArrivalDelay': 27.99}  
{'Airline': 'VX', 'avgArrivalDelay': 32.387}  
{'Airline': 'F9', 'avgArrivalDelay': 45.808}  
{'Airline': 'DL', 'avgArrivalDelay': 34.587}  
{'Airline': 'MQ', 'avgArrivalDelay': 40.959}  
{'Airline': 'OO', 'avgArrivalDelay': 35.497}  
{'Airline': 'HA', 'avgArrivalDelay': 14.775}  
{'Airline': 'NK', 'avgArrivalDelay': 37.987}  
{'Airline': 'B6', 'avgArrivalDelay': 41.441}  
{'Airline': 'EV', 'avgArrivalDelay': 34.486}
```



7. Days of months with respect to average of arrival delays. [Create a suitable plot using matplotlib/seaborn]

```
queryResult = list(  
    collection.aggregate([  
        {'$match': {'ARRIVAL_DELAY': {'$exists': True, '$ne':  
float('nan'), '$gte': 0}}},  
        {'$group': {'_id': '$DAY', 'avgArrivalDelay': {'$avg':  
'$ARRIVAL_DELAY'}}},  
        {'$sort': {'_id': 1}},  
        {'$project': {'_id': 0, 'DAY': '$_id', 'avgArrivalDelay': 1}}  
    ])  
)
```

```
for item in queryResult:  
    print(item)
```


```
df = pd.DataFrame(queryResult)
df.columns = ['AverageArrivalDelay', 'Day']

plt.figure(figsize=(15,5))
plt.xticks(ticks=df['Day'], rotation=0)
line_plot = sns.barplot(data=df, x='Day', y='AverageArrivalDelay')
```

```
plt.xlabel('Day_of_month')
plt.ylabel('Average Arrival Delay (in minutes)')
plt.title('Average Arrival Delay by each Day')
```

```
plt.show()
```

```
{'avgArrivalDelay': 41.932932932932935, 'DAY': 1}
{'avgArrivalDelay': 35.794701986754966, 'DAY': 2}
{'avgArrivalDelay': 39.91094890510949, 'DAY': 3}
{'avgArrivalDelay': 40.69406719085061, 'DAY': 4}
{'avgArrivalDelay': 40.506666666666667, 'DAY': 5}
{'avgArrivalDelay': 37.72260273972603, 'DAY': 6}
{'avgArrivalDelay': 29.276849642004773, 'DAY': 7}
{'avgArrivalDelay': 33.442996742671006, 'DAY': 8}
{'avgArrivalDelay': 29.541338582677167, 'DAY': 9}
{'avgArrivalDelay': 23.23248407643312, 'DAY': 10}
{'avgArrivalDelay': 31.445901639344264, 'DAY': 11}
{'avgArrivalDelay': 34.53033980582524, 'DAY': 12}
{'avgArrivalDelay': 26.832641770401107, 'DAY': 13}
{'avgArrivalDelay': 25.04553734061931, 'DAY': 14}
{'avgArrivalDelay': 25.240384615384617, 'DAY': 15}
{'avgArrivalDelay': 31.695106649937266, 'DAY': 16}
{'avgArrivalDelay': 35.7984126984127, 'DAY': 17}
{'avgArrivalDelay': 25.54357459379616, 'DAY': 18}
{'avgArrivalDelay': 23.172684458398745, 'DAY': 19}
{'avgArrivalDelay': 28.427272727272726, 'DAY': 20}
{'avgArrivalDelay': 33.12748344370861, 'DAY': 21}
{'avgArrivalDelay': 30.01761517615176, 'DAY': 22}
{'avgArrivalDelay': 25.006775067750677, 'DAY': 23}
{'avgArrivalDelay': 31.630872483221477, 'DAY': 24}
{'avgArrivalDelay': 31.50655737704918, 'DAY': 25}
{'avgArrivalDelay': 35.60516605166052, 'DAY': 26}
{'avgArrivalDelay': 34.065953654188945, 'DAY': 27}
{'avgArrivalDelay': 31.373358348968104, 'DAY': 28}
{'avgArrivalDelay': 23.3986013986014, 'DAY': 29}
{'avgArrivalDelay': 31.6355421686747, 'DAY': 30}
{'avgArrivalDelay': 29.12565445026178, 'DAY': 31}
```

8. Write a MongoDB aggregation pipeline to find the top 10 airports with the highest average total delay (DEPARTURE_DELAY + ARRIVAL_DELAY).

```
list(
  collection.aggregate([
    {
      '$match': {
        '$and': [
          {'ARRIVAL_DELAY': {'$exists': True, '$ne':
float('nan'), '$gte': 0}},
          {'DEPARTURE_DELAY': {'$exists': True, '$ne':
float('nan'), '$gte': 0}}
        ]
      },
    },
    {
      '$project': {
        'AIRPORT': '$ORIGIN_AIRPORT',
        'totalDelay': {
          '$add': ['$ARRIVAL_DELAY', '$DEPARTURE_DELAY']
        }
      }
    },
    {
      '$group': {'_id': '$AIRPORT', 'avgTotalDelay': {'$avg':
'$totalDelay'}}
    },
    {'$sort': {'avgTotalDelay': -1}},
    {'$limit': 10}
  ])
)

[{'_id': 'CDC', 'avgTotalDelay': 1082.0},
{'_id': 'PIH', 'avgTotalDelay': 964.0},
```

```
{ '_id': 'HOB', 'avgTotalDelay': 772.0},
{ '_id': 'ILG', 'avgTotalDelay': 551.0},
{ '_id': 'HIB', 'avgTotalDelay': 355.25},
{ '_id': 'SCE', 'avgTotalDelay': 335.0},
{ '_id': 'BRW', 'avgTotalDelay': 329.5},
{ '_id': 'DLH', 'avgTotalDelay': 312.0},
{ '_id': 'MBS', 'avgTotalDelay': 287.0},
{ '_id': 'PSG', 'avgTotalDelay': 281.0}]
```

9. Explain how you would create an index on the ORIGIN_AIRPORT and DESTINATION_AIRPORT fields to optimize queries filtering by these fields.

```
collection.create_index({'ORIGIN_AIRPORT':1, 'DESTINATION_AIRPORT':1})
'ORIGIN_AIRPORT_1_DESTINATION_AIRPORT_1'
```

10. Arrange weekdays with respect to the average arrival delays caused. [Create a suitable plot using matplotlib/seaborn]

```
queryResult = list(
    collection.aggregate([
        {'$match': {
            'ARRIVAL_DELAY': {'$exists': True, '$ne': float('nan')},
            '$gte': 0},
            'DAY_OF_WEEK': {'$gte': 1, '$lte': 5}
        }},
        {'$group': {'_id': '$DAY_OF_WEEK', 'avgArrivalDelay': {'$avg': '$ARRIVAL_DELAY'}}},
        {'$sort': {'avgArrivalDelay': -1}},
        {'$project': {'_id': 0, 'DAY_OF_WEEK': '$_id', 'avgArrivalDelay': 1}}
    ])
)

for item in queryResult:
    print(item)

df = pd.DataFrame(queryResult)
df.columns = ['AverageArrivalDelay', 'Day']
df['Day'] = df['Day'].map(lambda x: 'Monday' if x==1 else 'Tuesday' if x==2 else 'Wednesday' if x==3 else 'Thursday' if x==4 else 'Friday' if x==5 else 'Unknown')

plt.figure(figsize=(10,4))
sns.set_style("darkgrid")
line_plot = sns.lineplot(data=df, x='Day', y='AverageArrivalDelay',
```

```


marker='o')
# bar_plot = sns.barplot(data=df, x='Day', y='AverageArrivalDelay')

plt.xlabel('Weekdays')
plt.ylabel('Average Arrival Delay (in minutes)')
plt.title('Average Arrival Delay on each week day')

plt.show()

{'avgArrivalDelay': 34.818762781186095, 'DAY_OF_WEEK': 1}
{'avgArrivalDelay': 34.68610421836228, 'DAY_OF_WEEK': 2}
{'avgArrivalDelay': 31.2832689122032, 'DAY_OF_WEEK': 4}
{'avgArrivalDelay': 30.92716089349304, 'DAY_OF_WEEK': 3}
{'avgArrivalDelay': 29.700210194429847, 'DAY_OF_WEEK': 5}

```



11. Arrange Days of month as per cancellations done in descending order. [Create a suitable plot using matplotlib/seaborn]

```

queryResult = list(
    collection.aggregate([
        {'$group': {'_id': '$DAY', 'cancellations': {'$sum': '$CANCELLED'}}},
        # {'$sort': {'_id': 1}},
        {'$sort': {'cancellations': -1}},
        {'$project': {
            '_id': 0,
            'DAY': '_id',
            'cancellations': 1
        }}
    ])

```

```

)

for item in queryResult:
    print(item)


df = pd.DataFrame(queryResult)
df.columns = ['cancellations', 'DAY']
plt.figure(figsize=(10,4))
bar_plot = sns.barplot(data=df, x='DAY', y='cancellations')

plt.xlabel('Day of month')
plt.ylabel('Cancellations')
plt.title('Cancellations per day')

plt.show()

{'cancellations': 237, 'DAY': 1}
{'cancellations': 215, 'DAY': 5}
{'cancellations': 195, 'DAY': 2}
{'cancellations': 185, 'DAY': 27}
{'cancellations': 114, 'DAY': 26}
{'cancellations': 113, 'DAY': 4}
{'cancellations': 98, 'DAY': 28}
{'cancellations': 89, 'DAY': 9}
{'cancellations': 88, 'DAY': 3}
{'cancellations': 83, 'DAY': 15}
{'cancellations': 69, 'DAY': 23}
{'cancellations': 63, 'DAY': 16}
{'cancellations': 61, 'DAY': 25}
{'cancellations': 61, 'DAY': 8}
{'cancellations': 61, 'DAY': 21}
{'cancellations': 59, 'DAY': 17}
{'cancellations': 57, 'DAY': 24}
{'cancellations': 53, 'DAY': 6}
{'cancellations': 41, 'DAY': 22}
{'cancellations': 31, 'DAY': 7}
{'cancellations': 30, 'DAY': 14}
{'cancellations': 27, 'DAY': 10}
{'cancellations': 24, 'DAY': 11}
{'cancellations': 21, 'DAY': 20}
{'cancellations': 20, 'DAY': 12}
{'cancellations': 14, 'DAY': 18}
{'cancellations': 14, 'DAY': 13}
{'cancellations': 9, 'DAY': 19}
{'cancellations': 9, 'DAY': 31}
{'cancellations': 7, 'DAY': 30}
{'cancellations': 3, 'DAY': 29}

```



12. Find the busiest airports with respect to day of week. Represent it by using suitable plot.

```
queryResult = list(
    collection.aggregate([
        {'$group': {'_id':{'dayOfWeek':'$DAY_OF_WEEK',
'airport':'$ORIGIN_AIRPORT'}, 'totalFlights':{'$sum':1}}},
        {'$setWindowFields': {
            'partitionBy': '$_id.dayOfWeek',
            'sortBy': {'totalFlights':-1},
            'output': {'rank': {'$denseRank': {}}}}
        }},
        {'$match': {'rank':1}},
        {'$project': {'_id':0, 'Day':'$_id.dayOfWeek',
'Airport':'$_id.airport', 'totalFlights':1}}
    ])
)

for item in queryResult:
    print(item)


df = pd.DataFrame(queryResult)
df['Day'] = df['Day'].map(lambda x: 'Monday' if x==1 else 'Tuesday' if
x==2 else 'Wednesday' if x==3 else 'Thursday' if x==4 else 'Friday' if
x==5 else 'Saturday' if x==6 else 'Sunday' if x==7 else 'Unknown')
# df.columns = ['cancellations', 'DAY']
plt.figure(figsize=(10,4))
bar_plot = sns.barplot(data=df, x='Day', y='totalFlights',
hue='Airport')

plt.xlabel('Day')
```

```
plt.ylabel('TotalFlights')
plt.title('Busiest airport for each day of week')

plt.show()
```

```
{'totalFlights': 555, 'Day': 1, 'Airport': 'ATL'}
{'totalFlights': 475, 'Day': 2, 'Airport': 'ATL'}
{'totalFlights': 505, 'Day': 3, 'Airport': 'ATL'}
{'totalFlights': 556, 'Day': 4, 'Airport': 'ATL'}
{'totalFlights': 574, 'Day': 5, 'Airport': 'ATL'}
{'totalFlights': 413, 'Day': 6, 'Airport': 'ATL'}
{'totalFlights': 499, 'Day': 7, 'Airport': 'ATL'}
```



13. Find top 10 Airlines of US. Represent it by using suitable plot.

```
top_by_most_flights = list(
    collection.aggregate([
        {'$group': {'_id': '$AIRLINE', 'totalFlights': {'$sum': 1}}},
        {'$sort': {'totalFlights': -1}},
        {'$limit': 10}
    ])
)

top_by_least_avg_arrival_delay = list(
    collection.aggregate([
        {'$match': {'ARRIVAL_DELAY': {'$exists': True, '$ne':
float('nan'), '$gte': 0}}}},
        {'$group': {'_id': '$AIRLINE', 'avgArrivalDelay': {'$avg':
'$ARRIVAL_DELAY'}}},
        {'$sort': {'avgArrivalDelay': 1}},
        {'$limit': 10},
        {'$project': {'_id': 0, 'Airline': '$_id', 'avgArrivalDelay':
```

```

{'$round': ['$avgArrivalDelay',2]}}
    ])
)

top_by_least_avg_departure_delay = list(
    collection.aggregate([
        {'$match': {'DEPARTURE_DELAY': {'$exists': True, '$ne':
float('nan'), '$gte':0}}}},
        {'$group': {'_id': '$AIRLINE', 'avgDepartureDelay': {'$avg':
'$DEPARTURE_DELAY'}}},
        {'$sort': {'avgDepartureDelay': 1}},
        {'$limit': 10},
        {'$project': {'_id': 0, 'Airline': '$_id',
'avgDepartureDelay': {'$round': ['$avgDepartureDelay',2]}}}
    ])
)

top_by_least_avg_airline_delay = list(
    collection.aggregate([
        {'$match': {'AIRLINE_DELAY': {'$exists': True, '$ne':
float('nan'), '$gte':0}}}},
        {'$group': {'_id': '$AIRLINE', 'avgAirlineDelay': {'$avg':
'$AIRLINE_DELAY'}}},
        {'$sort': {'avgAirlineDelay': 1}},
        {'$limit': 10},
        {'$project': {'_id': 0, 'Airline': '$_id', 'avgAirlineDelay':
{'$round': ['$avgAirlineDelay',2]}}}
    ])
)

# convert all data into DataFrame format
df_flights = pd.DataFrame(top_by_most_flights)
df_flights.columns = ['Airline', 'Total_Flights']

df_arrival_delay = pd.DataFrame(top_by_least_avg_arrival_delay)
df_arrival_delay.columns = ['Airline', 'Avg_Arrival_Delay']

df_departure_delay = pd.DataFrame(top_by_least_avg_departure_delay)
df_departure_delay.columns = ['Airline', 'Avg_Departure_Delay']

df_airline_delay = pd.DataFrame(top_by_least_avg_airline_delay)
df_airline_delay.columns = ['Airline', 'Avg_Airline_Delay']

# subplots
fig, axes = plt.subplots(2, 2, figsize=(18, 12))

# Top 10 Airlines by Total Number of Flights
sns.barplot(data=df_flights, x='Airline', y='Total_Flights',

```

```

palette='viridis', ax=axes[0, 0])
axes[0,0].set_title('Top 10 Airlines by Total Number of Flights')
axes[0,0].set_xlabel('Airline')
axes[0,0].set_ylabel('Total Flights')
axes[0,0].tick_params(axis='x', rotation=45)

# Top 10 Airlines by Least Average Arrival Delay
sns.barplot(data=df_arrival_delay, x='Airline', y='Avg_Arrival_Delay',
palette='coolwarm', ax=axes[0, 1])
axes[0,1].set_title('Top 10 Airlines by Least Average Arrival Delay')
axes[0,1].set_xlabel('Airline')
axes[0,1].set_ylabel('Average Arrival Delay (Minutes)')
axes[0,1].tick_params(axis='x', rotation=45)

# Top 10 Airlines by Least Average Departure Delay
sns.barplot(data=df_departure_delay, x='Airline',
y='Avg_Departure_Delay', palette='coolwarm', ax=axes[1, 0])
axes[1,0].set_title('Top 10 Airlines by Least Average Departure
Delay')
axes[1,0].set_xlabel('Airline')
axes[1,0].set_ylabel('Average Departure Delay (Minutes)')
axes[1,0].tick_params(axis='x', rotation=45)

# Top 10 airlines by Least Average Airline Delay
sns.barplot(data=df_airline_delay, x='Airline', y='Avg_Airline_Delay',
palette='coolwarm', ax=axes[1, 1])
axes[1,1].set_title('Top 10 Airlines by Least Average Airline Delay')
axes[1,1].set_xlabel('Airline')
axes[1,1].set_ylabel('Average Airline Delay (Minutes)')
axes[1,1].tick_params(axis='x', rotation=45)

plt.show()

C:\Users\Administrator\AppData\Local\Temp\
ipykernel_18224\4284932834.py:58: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.

sns.barplot(data=df_flights, x='Airline', y='Total_Flights',
palette='viridis', ax=axes[0, 0])
C:\Users\Administrator\AppData\Local\Temp\
ipykernel_18224\4284932834.py:65: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.

sns.barplot(data=df_arrival_delay, x='Airline',

```







```
y='Avg_Arrival_Delay', palette='coolwarm', ax=axes[0, 1])
C:\Users\Administrator\AppData\Local\Temp\
ipykernel_18224\4284932834.py:72: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(data=df_departure_delay, x='Airline',
y='Avg_Departure_Delay', palette='coolwarm', ax=axes[1, 0])
C:\Users\Administrator\AppData\Local\Temp\
ipykernel_18224\4284932834.py:79: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(data=df_airline_delay, x='Airline',
y='Avg_Airline_Delay', palette='coolwarm', ax=axes[1, 1])
```



14. Finding airlines that make the maximum, minimum number of cancellations.

```
list(
    collection.aggregate([
        {'$group': {'_id': '$AIRLINE', 'totalCancellations': {'$sum': '$CANCELLED'}}},
        {'$facet': {
            'maxCancellations': [{'$sort': {'totalCancellations': -1}}, {'$limit': 1}],
            'minCancellations': [{'$sort': {'totalCancellations': 1}}, {'$limit': 1}]
        }}
    ])
)
```

```
[{'maxCancellations': [{'_id': 'MQ', 'totalCancellations': 414}],
 'minCancellations': [{'_id': 'HA', 'totalCancellations': 3}]]
```

15. Find and show airlines names in descending that make the most number of diversions made. [Create a suitable plot using matplotlib/seaborn]

```
queryResult = list(
    collection.aggregate([
        {'$match': {'DIVERTED': {'$eq': 1}}},
        {'$group': {'_id': '$AIRLINE', 'totalDiversions': {'$sum': 1}}},
        {'$sort': {'totalDiversions': -1}},
        {'$project': {'_id': 0, 'Airline': '$_id', 'totalDiversions': 1}}
    ])
)
```

```
for item in queryResult:
    print(item)
```

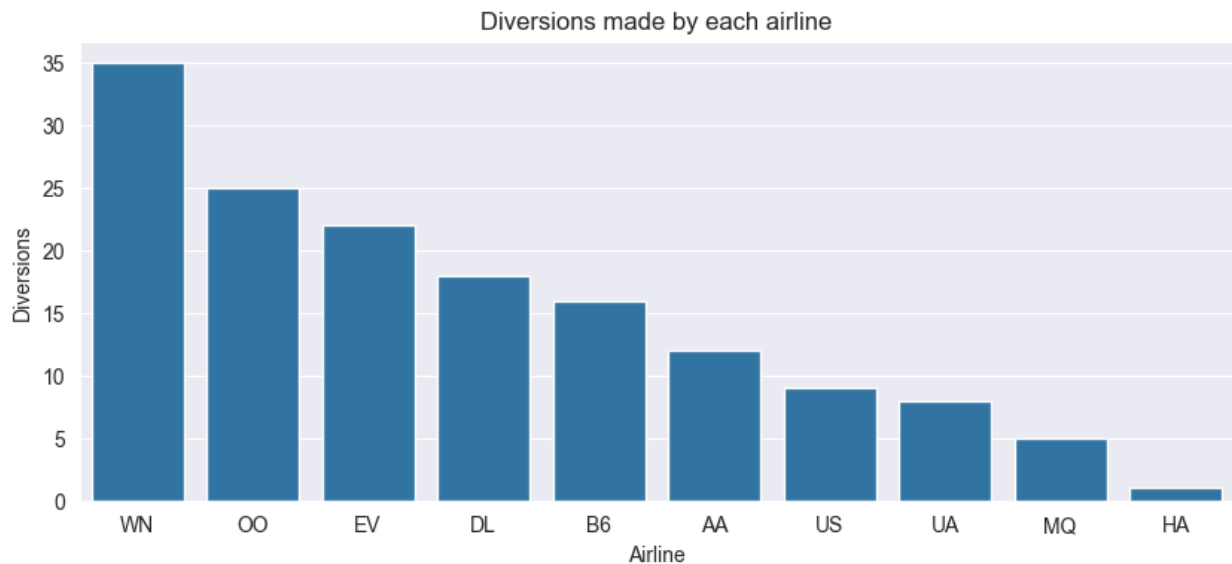
```
df = pd.DataFrame(queryResult)
df.columns = ['diversions', 'airline']
plt.figure(figsize=(10,4))
bar_plot = sns.barplot(data=df, x='airline', y='diversions')
```

```
plt.xlabel('Airline')
plt.ylabel('Diversions')
plt.title('Diversions made by each airline')
```

```
plt.show()
```

```
{ 'totalDiversions': 35, 'Airline': 'WN' }
{ 'totalDiversions': 25, 'Airline': '00' }
```

```
{'totalDiversions': 22, 'Airline': 'EV'}
{'totalDiversions': 18, 'Airline': 'DL'}
{'totalDiversions': 16, 'Airline': 'B6'}
{'totalDiversions': 12, 'Airline': 'AA'}
{'totalDiversions': 9, 'Airline': 'US'}
{'totalDiversions': 8, 'Airline': 'UA'}
{'totalDiversions': 5, 'Airline': 'MQ'}
{'totalDiversions': 1, 'Airline': 'HA'}
```



16. Finding days of month that see the most number of diversion and delays.

```
list(
  collection.aggregate([
    {'$group': {
      '_id': '$DAY',
      'totalDiversions': {'$sum': '$DIVERTED'},
      'totalDelayCount': {
        '$sum': {
          '$cond': [{
            '$or': [
              {'$gt': ['$ARRIVAL_DELAY', 0]}, {'$gt':
['$DEPARTURE_DELAY', 0]}
            ]
          }, 1, 0]
        }
      }
    }},
    {'$sort': {'totalDiversions': -1, 'totalDelayCount': -1}},
    {'$project': {'_id': 0, 'Day': '$_id', 'totalDiversions': 1,
'totalDelayCount': 1}}
```

```

    })
  )
  [
    {'totalDiversions': 15, 'totalDelayCount': 1573, 'Day': 2},
    {'totalDiversions': 13, 'totalDelayCount': 1206, 'Day': 1},
    {'totalDiversions': 12, 'totalDelayCount': 1582, 'Day': 4},
    {'totalDiversions': 11, 'totalDelayCount': 1500, 'Day': 5},
    {'totalDiversions': 9, 'totalDelayCount': 1231, 'Day': 9},
    {'totalDiversions': 8, 'totalDelayCount': 659, 'Day': 14},
    {'totalDiversions': 7, 'totalDelayCount': 1418, 'Day': 6},
    {'totalDiversions': 6, 'totalDelayCount': 1063, 'Day': 7},
    {'totalDiversions': 6, 'totalDelayCount': 845, 'Day': 23},
    {'totalDiversions': 5, 'totalDelayCount': 1511, 'Day': 3},
    {'totalDiversions': 5, 'totalDelayCount': 1123, 'Day': 8},
    {'totalDiversions': 5, 'totalDelayCount': 814, 'Day': 18},
    {'totalDiversions': 5, 'totalDelayCount': 734, 'Day': 11},
    {'totalDiversions': 5, 'totalDelayCount': 391, 'Day': 30},
    {'totalDiversions': 4, 'totalDelayCount': 966, 'Day': 12},
    {'totalDiversions': 4, 'totalDelayCount': 950, 'Day': 16},
    {'totalDiversions': 4, 'totalDelayCount': 817, 'Day': 20},
    {'totalDiversions': 4, 'totalDelayCount': 698, 'Day': 21},
    {'totalDiversions': 4, 'totalDelayCount': 645, 'Day': 28},
    {'totalDiversions': 3, 'totalDelayCount': 917, 'Day': 26},
    {'totalDiversions': 3, 'totalDelayCount': 762, 'Day': 17},
    {'totalDiversions': 3, 'totalDelayCount': 689, 'Day': 27},
    {'totalDiversions': 3, 'totalDelayCount': 242, 'Day': 31},
    {'totalDiversions': 2, 'totalDelayCount': 876, 'Day': 13},
    {'totalDiversions': 1, 'totalDelayCount': 800, 'Day': 19},
    {'totalDiversions': 1, 'totalDelayCount': 767, 'Day': 15},
    {'totalDiversions': 1, 'totalDelayCount': 762, 'Day': 10},
    {'totalDiversions': 1, 'totalDelayCount': 740, 'Day': 25},
    {'totalDiversions': 1, 'totalDelayCount': 336, 'Day': 29},
    {'totalDiversions': 0, 'totalDelayCount': 863, 'Day': 22},
    {'totalDiversions': 0, 'totalDelayCount': 694, 'Day': 24}]

```

17. Write a MongoDB query to find the flights with the shortest and longest AIR_TIME. Return the flightNumber, airline, and AIR_TIME.

```

list(
  collection.aggregate([
    {'$match': {'AIR_TIME': {'$exists': True, '$ne':
float('nan')}}},
    {'$project': {'_id': 0, 'FLIGHT_NUMBER': 1, 'AIRLINE': 1,
'AIR_TIME': 1}},
    {'$facet': {
      'maxAirTime': [{'$sort': {'AIR_TIME': -1}}, {'$limit': 1}],
      'minAirTime': [{'$sort': {'AIR_TIME': 1}}, {'$limit': 1}]
    }}
  ])

```

```

    ])
)

[{'maxAirTime': [{'AIRLINE': 'UA', 'FLIGHT_NUMBER': 15, 'AIR_TIME': 654.0}],
  'minAirTime': [{'AIRLINE': 'AS', 'FLIGHT_NUMBER': 65, 'AIR_TIME': 9.0}]]

```

18. Finding all diverted Route from a source to destination Airport & which route is the most diverted route.

```

list(
  collection.aggregate([
    {'$group': {'_id': {'origin': '$ORIGIN_AIRPORT',
'destination': '$DESTINATION_AIRPORT'}, 'totalDiverted':
{'$sum': '$DIVERTED'}}},
    # {'$sort': {'totalDiverted': -1}},
    {'$setWindowFields': {
      'sortBy': {'totalDiverted': -1},
      'output': {'rank': {'$denseRank': {}}}}
    }},
    {'$match': {'rank': 1}},
    {'$project': {'_id': 0, 'origin': '$_id.origin',
'destination': '$_id.destination', 'totalDiverted': 1}},
  ])
)

[{'totalDiverted': 2, 'origin': 'PHL', 'destination': 'SAN'},
{'totalDiverted': 2, 'origin': 'IAH', 'destination': 'ASE'},
{'totalDiverted': 2, 'origin': 'HOU', 'destination': 'DAL'},
{'totalDiverted': 2, 'origin': 'STT', 'destination': 'PHL'},
{'totalDiverted': 2, 'origin': 'TPA', 'destination': 'LGA'},
{'totalDiverted': 2, 'origin': 'JFK', 'destination': 'SEA'},
{'totalDiverted': 2, 'origin': 'JFK', 'destination': 'EGE'},
{'totalDiverted': 2, 'origin': 'CLT', 'destination': 'IAH'},
{'totalDiverted': 2, 'origin': 'ORD', 'destination': 'ASE'}]

```

19. Write a MongoDB aggregation pipeline to calculate the all aggregated values for departure delay (DEPARTURE_DELAY) and arrival delay (ARRIVAL_DELAY) for each airline, excluding flights that were either cancelled or diverted.

```

list(
  collection.aggregate([
    {
      '$match': {
        '$and': [
          {'ARRIVAL_DELAY': {'$exists': True, '$ne':

```

```

float('nan'))}},
float('nan'))}},
    {'DEPARTURE_DELAY': {'$exists': True, '$ne':
    {'DIVERTED': 0},
    {'CANCELLED': 0}
    ]
    },
    {
        '$group': {
            '_id': '$AIRLINE',
            'minArrivalDelay': {'$min': '$ARRIVAL_DELAY'},
            'maxArrivalDelay': {'$max': '$ARRIVAL_DELAY'},
            'avgArrivalDelay': {'$avg': '$ARRIVAL_DELAY'},
            'medianArrivalDelay': {'$median':
{'input': '$ARRIVAL_DELAY', 'method': 'approximate'}},
            'minDepartureDelay': {'$min': '$DEPARTURE_DELAY'},
            'maxDepartureDelay': {'$max': '$DEPARTURE_DELAY'},
            'avgDepartureDelay': {'$avg': '$DEPARTURE_DELAY'},
            'medianDepartureDelay': {'$median':
{'input': '$DEPARTURE_DELAY', 'method': 'approximate'}}
        },
        {
            '$project': {
                '_id': 0,
                'airline': '$_id',
                'minArrivalDelay': 1,
                'maxArrivalDelay': 1,
                'avgArrivalDelay': {'$round': ['$avgArrivalDelay',
2]}},
                'medianArrivalDelay': 1,
                'minDepartureDelay': 1,
                'maxDepartureDelay': 1,
                'avgDepartureDelay': {'$round': ['$avgDepartureDelay',
2]}},
                'medianDepartureDelay': 1
            }
        }
    })
)

[{'minArrivalDelay': -76.0,
'maxArrivalDelay': 522.0,
'medianArrivalDelay': -1.0,
'minDepartureDelay': -24.0,
'maxDepartureDelay': 468.0,
'medianDepartureDelay': -1.0,
'airline': 'B6',
'avgArrivalDelay': 13.96,

```

```
'avgDepartureDelay': 15.89},
{'minArrivalDelay': -50.0,
 'maxArrivalDelay': 571.0,
 'medianArrivalDelay': -1.0,
 'minDepartureDelay': -24.0,
 'maxDepartureDelay': 526.0,
 'medianDepartureDelay': -2.0,
 'airline': 'EV',
 'avgArrivalDelay': 10.88,
 'avgDepartureDelay': 11.42},
{'minArrivalDelay': -45.0,
 'maxArrivalDelay': 511.0,
 'medianArrivalDelay': 3.0,
 'minDepartureDelay': -25.0,
 'maxDepartureDelay': 494.0,
 'medianDepartureDelay': 0.0,
 'airline': 'MQ',
 'avgArrivalDelay': 19.23,
 'avgDepartureDelay': 16.76},
{'minArrivalDelay': -50.0,
 'maxArrivalDelay': 542.0,
 'medianArrivalDelay': -2.0,
 'minDepartureDelay': -36.0,
 'maxDepartureDelay': 540.0,
 'medianDepartureDelay': -3.0,
 'airline': '00',
 'avgArrivalDelay': 10.15,
 'avgDepartureDelay': 11.47},
{'minArrivalDelay': -73.0,
 'maxArrivalDelay': 521.0,
 'medianArrivalDelay': -4.0,
 'minDepartureDelay': -15.0,
 'maxDepartureDelay': 490.0,
 'medianDepartureDelay': 0.0,
 'airline': 'WN',
 'avgArrivalDelay': 3.7,
 'avgDepartureDelay': 10.06},
{'minArrivalDelay': -61.0,
 'maxArrivalDelay': 407.0,
 'medianArrivalDelay': -7.0,
 'minDepartureDelay': -42.0,
 'maxDepartureDelay': 400.0,
 'medianDepartureDelay': -4.0,
 'airline': 'AS',
 'avgArrivalDelay': -1.53,
 'avgDepartureDelay': 2.31},
{'minArrivalDelay': -52.0,
 'maxArrivalDelay': 736.0,
 'medianArrivalDelay': 0.0,
```

```
'minDepartureDelay': -17.0,
'maxDepartureDelay': 715.0,
'medianDepartureDelay': -4.0,
'airline': 'HA',
'avgArrivalDelay': 4.07,
'avgDepartureDelay': 1.19},
{'minArrivalDelay': -64.0,
'maxArrivalDelay': 328.0,
'medianArrivalDelay': -3.0,
'minDepartureDelay': -21.0,
'maxDepartureDelay': 327.0,
'medianDepartureDelay': -2.0,
'airline': 'US',
'avgArrivalDelay': 5.98,
'avgDepartureDelay': 7.75},
{'minArrivalDelay': -65.0,
'maxArrivalDelay': 323.0,
'medianArrivalDelay': -4.0,
'minDepartureDelay': -15.0,
'maxDepartureDelay': 309.0,
'medianDepartureDelay': -1.0,
'airline': 'VX',
'avgArrivalDelay': 5.13,
'avgDepartureDelay': 9.86},
{'minArrivalDelay': -39.0,
'maxArrivalDelay': 523.0,
'medianArrivalDelay': 0.0,
'minDepartureDelay': -22.0,
'maxDepartureDelay': 546.0,
'medianDepartureDelay': -1.0,
'airline': 'NK',
'avgArrivalDelay': 14.21,
'avgDepartureDelay': 15.53},
{'minArrivalDelay': -38.0,
'maxArrivalDelay': 505.0,
'medianArrivalDelay': 6.0,
'minDepartureDelay': -32.0,
'maxDepartureDelay': 499.0,
'medianDepartureDelay': 2.0,
'airline': 'F9',
'avgArrivalDelay': 24.1,
'avgDepartureDelay': 23.51},
{'minArrivalDelay': -62.0,
'maxArrivalDelay': 1170.0,
'medianArrivalDelay': -7.0,
'minDepartureDelay': -26.0,
'maxDepartureDelay': 1166.0,
'medianDepartureDelay': -1.0,
'airline': 'DL',
```



```

    'avgArrivalDelay': 2.81,
    'avgDepartureDelay': 9.92},
    {'minArrivalDelay': -56.0,
     'maxArrivalDelay': 470.0,
     'medianArrivalDelay': -3.0,
     'minDepartureDelay': -23.0,
     'maxDepartureDelay': 473.0,
     'medianDepartureDelay': 2.0,
     'airline': 'UA',
     'avgArrivalDelay': 6.7,
     'avgDepartureDelay': 14.17},
    {'minArrivalDelay': -62.0,
     'maxArrivalDelay': 1295.0,
     'medianArrivalDelay': -3.0,
     'minDepartureDelay': -23.0,
     'maxDepartureDelay': 1264.0,
     'medianDepartureDelay': -2.0,
     'airline': 'AA',
     'avgArrivalDelay': 8.39,
     'avgDepartureDelay': 11.41}]

```

20. Write a MongoDB query to find all flights that were delayed due to WEATHER_DELAY but were not cancelled or diverted. Include the flightNumber, airline, originAirport, and destinationAirport in the results.

```

list(
  collection.find({
    'WEATHER_DELAY': {'$exists': True, '$ne': float('nan')},
    '$gt': 0},
    {'CANCELLED': {'$ne': 1},
     'DIVERTED': {'$ne': 1}}
  }, {
    '_id': 0, 'FLIGHT_NUMBER': 1, 'AIRLINE': 1,
    'ORIGIN_AIRPORT': 1, 'DESTINATION_AIRPORT': 1
  })
).limit(10)

[{'AIRLINE': 'UA',
 'FLIGHT_NUMBER': 532,
 'ORIGIN_AIRPORT': 'ORD',
 'DESTINATION_AIRPORT': 'DCA'},
 {'AIRLINE': 'US',
 'FLIGHT_NUMBER': 1784,
 'ORIGIN_AIRPORT': 'BWI',
 'DESTINATION_AIRPORT': 'PHX'},
 {'AIRLINE': 'MQ',

```

```

    'FLIGHT_NUMBER': 3019,
    'ORIGIN_AIRPORT': 'ORD',
    'DESTINATION_AIRPORT': 'OKC'},
  {'AIRLINE': 'MQ',
    'FLIGHT_NUMBER': 3564,
    'ORIGIN_AIRPORT': 'GSO',
    'DESTINATION_AIRPORT': 'LGA'},
  {'AIRLINE': 'UA',
    'FLIGHT_NUMBER': 1667,
    'ORIGIN_AIRPORT': 'ORD',
    'DESTINATION_AIRPORT': 'PDX'},
  {'AIRLINE': 'DL',
    'FLIGHT_NUMBER': 1788,
    'ORIGIN_AIRPORT': 'ATL',
    'DESTINATION_AIRPORT': 'MEM'},
  {'AIRLINE': 'DL',
    'FLIGHT_NUMBER': 424,
    'ORIGIN_AIRPORT': 'JFK',
    'DESTINATION_AIRPORT': 'LAX'},
  {'AIRLINE': 'MQ',
    'FLIGHT_NUMBER': 3201,
    'ORIGIN_AIRPORT': 'ORD',
    'DESTINATION_AIRPORT': 'BNA'},
  {'AIRLINE': 'UA',
    'FLIGHT_NUMBER': 1718,
    'ORIGIN_AIRPORT': 'LAX',
    'DESTINATION_AIRPORT': 'KOA'},
  {'AIRLINE': 'DL',
    'FLIGHT_NUMBER': 338,
    'ORIGIN_AIRPORT': 'DTW',
    'DESTINATION_AIRPORT': 'ATL'}]

```

21. Write a MongoDB query to find all flights that were delayed both at departure (DEPARTURE_DELAY) and arrival (ARRIVAL_DELAY). Return the count of such Flights which are delayed.

```

list(
  collection.aggregate([
    {'$match': {'DEPARTURE_DELAY': {'$gt': 0}, 'ARRIVAL_DELAY':
{'$gt': 0}}},
    {'$count': 'countDelayedFlights'}
  ])
)
[{'countDelayedFlights': 16481}]

```

22. Write a MongoDB query to calculate the frequency of flight takeoffs and landings within defined time intervals (e.g., every hour) throughout the day. Generate a Suitable Plot.

```
queryResult = list(
    collection.aggregate([
        {'$project': {
            'departureHour': {'$floor': {'$divide':
['$DEPARTURE_TIME', 100]}},
            'arrivalHour': {'$floor': {'$divide': ['$ARRIVAL_TIME',
100]}}}
        }},
        {
            '$facet': {
                'arrivalCounts': [
                    {'$group': {'_id': '$arrivalHour', 'count':
{'$sum': 1}}},
                    {'$sort': {'_id': 1}}
                ],
                'departureCounts': [
                    {'$group': {'_id': '$departureHour', 'count':
{'$sum': 1}}},
                    {'$sort': {'_id': 1}}
                ]
            }
        }
    ])
)

arrival_df = pd.DataFrame(queryResult[0]['arrivalCounts'])
departure_df = pd.DataFrame(queryResult[0]['departureCounts'])

arrival_df.columns = ['Hour', 'ArrivalCount']
arrival_df = arrival_df.dropna()      # some Hour values were NaN, so
dropping them
departure_df.columns = ['Hour', 'DepartureCount']
departure_df = departure_df.dropna()
temp_df = arrival_df.merge(departure_df)
print(temp_df)

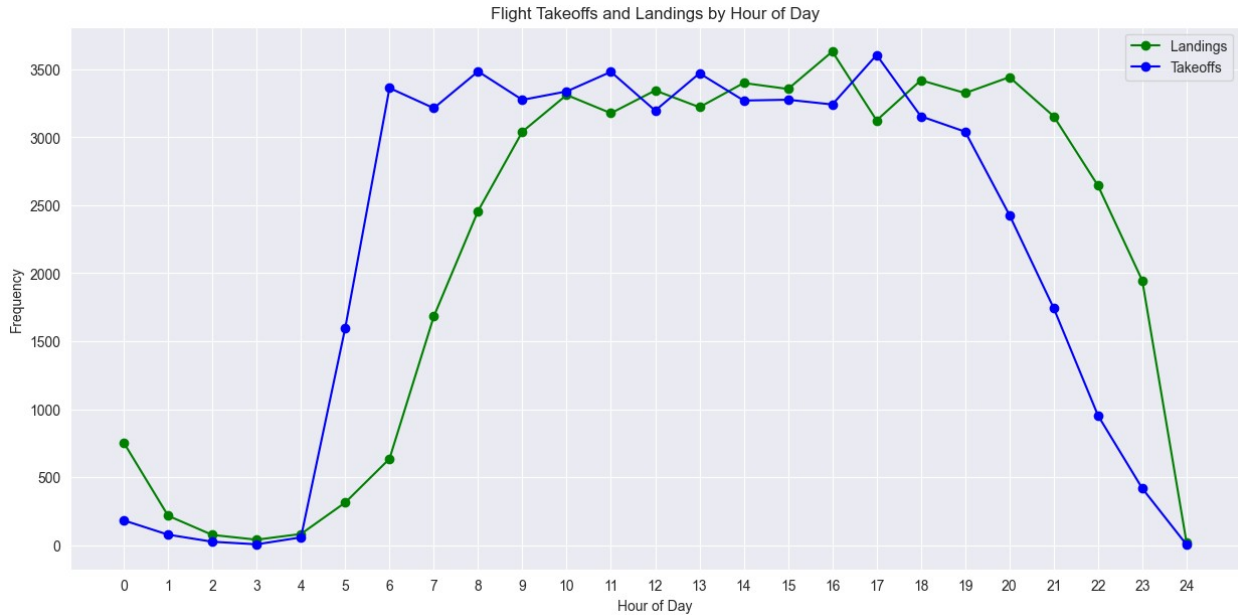
plt.figure(figsize=(15, 7))
plt.plot(arrival_df['Hour'], arrival_df['ArrivalCount'], marker='o',
color='green', label='Landings')
plt.plot(departure_df['Hour'], departure_df['DepartureCount'],
marker='o', color='blue', label='Takeoffs')

plt.xlabel('Hour of Day')
```

```
plt.ylabel('Frequency')
plt.title('Flight Takeoffs and Landings by Hour of Day')
plt.legend()
sns.set_style("darkgrid")
plt.xticks(ticks=arrival_df['Hour'], rotation=0)

plt.show()
```

	Hour	ArrivalCount	DepartureCount
0	0.0	754	183
1	1.0	217	78
2	2.0	76	26
3	3.0	41	7
4	4.0	83	58
5	5.0	314	1598
6	6.0	636	3362
7	7.0	1682	3214
8	8.0	2458	3484
9	9.0	3041	3275
10	10.0	3311	3337
11	11.0	3177	3481
12	12.0	3344	3196
13	13.0	3221	3470
14	14.0	3399	3270
15	15.0	3355	3276
16	16.0	3630	3241
17	17.0	3124	3605
18	18.0	3419	3154
19	19.0	3325	3041
20	20.0	3442	2428
21	21.0	3154	1743
22	22.0	2644	954
23	23.0	1946	418
24	24.0	18	4



23. Write a MongoDB query to calculate the frequency of flight takeoffs and landings within defined week of day. Generate a Sutable Plot.

```
queryResult = list(
    collection.aggregate([
        {'$group': {
            '_id': '$DAY_OF_WEEK',
            'takeoffCount': {'$sum': {'$cond': [{'$gte':
['$DEPARTURE_TIME', 0]}, 1, 0]}},
            'landingCount': {'$sum': {'$cond': [{'$gte':
['$ARRIVAL_TIME', 0]}, 1, 0]}},
            }},
        {'$sort': {'_id': 1}}
    ])
)


df = pd.DataFrame(queryResult)
print(df)
df.columns = ['DAY', 'Take_Off_Count', 'Landing_Count']
df['DAY'] = df['DAY'].map(lambda x: 'Monday' if x==1 else 'Tuesday' if
x==2 else 'Wednesday' if x==3 else 'Thursday' if x==4 else 'Friday' if
x==5 else 'Saturday' if x==6 else 'Sunday' if x==7 else 'Unknown')

df_melted = df.melt(id_vars='DAY', value_vars=['Take_Off_Count',
'Landing_Count'], var_name='Type', value_name='Count')
# print(df_melted)
plt.figure(figsize=(12, 6))
sns.barplot(data=df_melted, x='DAY', y='Count', hue='Type')
```

```
plt.xlabel('Day of Week')
plt.ylabel('Frequency')
plt.title('Takeoff and Landing Frequency Per Day')
sns.set_style("darkgrid")
plt.legend()

plt.show()
```

	_id	takeoffCount	landingCount
0	1	8047	8024
1	2	7423	7414
2	3	7452	7441
3	4	8098	8084
4	5	8756	8744
5	6	6705	6692
6	7	7422	7412



24. Write a MongoDB query to find all flights that departed between 6 AM and 12 PM (noon) local time, regardless of the date. Return the flightNumber, airline, and departureTime. Generate a Bar Plot using Time (x-axis) and Frequency (y-axis).

```
queryResult = list(
    collection.aggregate([
        {'$match': {'DEPARTURE_TIME': {'$gte':600, '$lte':1200}}},
        {'$project': {'_id':0, 'FLIGHT_NUMBER':1, 'AIRLINE':1,
            'DEPARTURE_TIME':1}},
```

```

    ])
)

df = pd.DataFrame(queryResult)
print(df)
df['DEPARTURE_HOUR'] = df['DEPARTURE_TIME']//100

plt.figure(figsize=(10,4))
histplot = sns.histplot(data=df['DEPARTURE_HOUR'], discrete=True,
bins=range(6, 13))

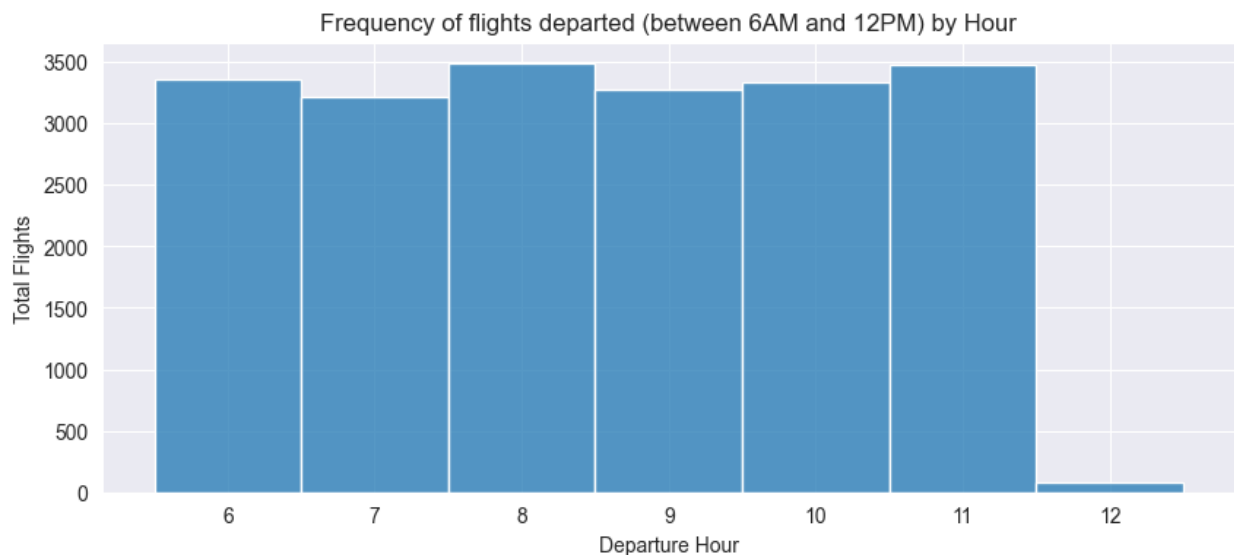
plt.xlabel('Departure Hour')
plt.ylabel('Total Flights')
plt.title('Frequency of flights departed (between 6AM and 12PM) by
Hour')

plt.show()

```

	AIRLINE	FLIGHT_NUMBER	DEPARTURE_TIME
0	EV	5170	954.0
1	00	5166	924.0
2	UA	321	947.0
3	WN	2685	633.0
4	MQ	3196	1051.0
...
20225	AA	1349	627.0
20226	B6	1567	1029.0
20227	AA	1113	857.0
20228	US	661	621.0
20229	DL	1318	1003.0

[20230 rows x 3 columns]



25. When is the best time of day/day of week/time of a year to fly with minimum delays?

```
# finding best time of day to fly
```

```
queryResult = list(
    collection.aggregate([
        {'$match': {
            'ARRIVAL_DELAY': {'$exists': True, '$ne': float('nan')},
            '$gte': 0},
            'DEPARTURE_DELAY': {'$exists': True, '$ne': float('nan')},
            '$gte': 0}
        }},
        {'$project': {
            '_id': 0,
            'hourOfDay': {'$floor': {'$divide': ['$DEPARTURE_TIME',
100]}},
            'totalDelay': {'$add': ['$ARRIVAL_DELAY',
'$DEPARTURE_DELAY']}
        }},
        {'$group': {'_id': '$hourOfDay', 'avgDelay': {'$avg':
'$totalDelay'}}},
        {'$sort': {'_id': 1}}
    ])
)
```

```
for record in queryResult:
    print(record)
```

```
df = pd.DataFrame(queryResult)
df.columns = ['Hour', 'Average_Delay_Per_Hour']
plt.figure(figsize=(10,5))
bar_plot = sns.barplot(data=df, x='Hour', y='Average_Delay_Per_Hour')
```


```
plt.xlabel('Hour of Day')
plt.ylabel('Average Delay per Hour')
plt.show()
```

```
# insight: least delay is arouun 6AM morning, with average total delay
of 43.8 minutes
```

```
{'_id': 0.0, 'avgDelay': 220.68}
{'_id': 1.0, 'avgDelay': 240.5121951219512}
{'_id': 2.0, 'avgDelay': 291.05882352941177}
{'_id': 3.0, 'avgDelay': 75.0}
{'_id': 4.0, 'avgDelay': 385.2}
{'_id': 5.0, 'avgDelay': 49.93421052631579}
{'_id': 6.0, 'avgDelay': 43.801781737193764}
{'_id': 7.0, 'avgDelay': 44.61577181208054}
{'_id': 8.0, 'avgDelay': 56.35}
{'_id': 9.0, 'avgDelay': 67.21552723059096}
```



```
{'_id': 10.0, 'avgDelay': 68.26398429833169}
{'_id': 11.0, 'avgDelay': 71.37534246575342}
{'_id': 12.0, 'avgDelay': 75.20092592592593}
{'_id': 13.0, 'avgDelay': 72.85085574572128}
{'_id': 14.0, 'avgDelay': 75.33279352226721}
{'_id': 15.0, 'avgDelay': 79.15071770334929}
{'_id': 16.0, 'avgDelay': 70.90981012658227}
{'_id': 17.0, 'avgDelay': 77.71942446043165}
{'_id': 18.0, 'avgDelay': 83.96085672082718}
{'_id': 19.0, 'avgDelay': 90.75827559661278}
{'_id': 20.0, 'avgDelay': 98.62594776748105}
{'_id': 21.0, 'avgDelay': 104.75267538644471}
{'_id': 22.0, 'avgDelay': 125.50766283524904}
{'_id': 23.0, 'avgDelay': 189.445}
{'_id': 24.0, 'avgDelay': 203.0}
```



```
# finding best day of week to fly

queryResult = list(
    collection.aggregate([
        {'$match': {
            'ARRIVAL_DELAY': {'$exists': True, '$ne': float('nan')},
            '$gte': 0},
            'DEPARTURE_DELAY': {'$exists': True, '$ne': float('nan')},
            '$gte': 0}
        }},
        {'$group': {'_id': '$DAY_OF_WEEK', 'avgDelay': {'$avg':
{'$add': ['$ARRIVAL_DELAY', '$DEPARTURE_DELAY']}}}},
        {'$sort': {'_id': 1}},
```

```

        {'$project': {
            '_id': 0,
            'DayOfWeek': '$_id',
            'avgDelay': 1
        }},
    ]
)

for record in queryResult:
    print(record)


df = pd.DataFrame(queryResult)
df.columns = ['Average_Delay_of_Day', 'Day_Of_Week']
df['Day_Of_Week'] = df['Day_Of_Week'].map(lambda x: 'Monday' if x==1
else 'Tuesday' if x==2 else 'Wednesday' if x==3 else 'Thursday' if
x==4 else 'Friday' if x==5 else 'Saturday' if x==6 else 'Sunday' if
x==7 else 'Unknown')

plt.figure(figsize=(10,5))
bar_plot = sns.barplot(data=df, x='Day_Of_Week',
y='Average_Delay_of_Day')

plt.xlabel('Day of Week')
plt.ylabel('Average Delay per Day')
plt.show()
# insight: day of week with least delay is Friday, with average delay
of 72.8 minutes

{'avgDelay': 83.59503592423253, 'DayOfWeek': 1}
{'avgDelay': 85.65494137353434, 'DayOfWeek': 2}
{'avgDelay': 78.26551880380607, 'DayOfWeek': 3}
{'avgDelay': 76.52054794520548, 'DayOfWeek': 4}
{'avgDelay': 72.89926803764378, 'DayOfWeek': 5}
{'avgDelay': 79.11367050575863, 'DayOfWeek': 6}
{'avgDelay': 87.29958599924727, 'DayOfWeek': 7}

```



```
# finding best time of year to fly
```

```
queryResult = list(
    collection.aggregate([
        {'$match': {
            'ARRIVAL_DELAY': {'$exists': True, '$ne': float('nan')},
            'DEPARTURE_DELAY': {'$exists': True, '$ne': float('nan')}
        }},
        {'$group': {
            '_id': '$MONTH',
            'avgDelay': {'$avg': {'$add': ['$ARRIVAL_DELAY',
'$DEPARTURE_DELAY']}}
        }},
        {'$sort': {'_id': 1}}
    ])
)
```


```
for record in queryResult:
    print(record)
```

```
df = pd.DataFrame(queryResult)
df.columns = ['Month', 'Average_Delay_of_Month']
df['Month'] = df['Month'].map(lambda x: 'January' if x==1 else
'February' if x==2 else 'March' if x==3 else 'April' if x==4 else
'May' if x==5 else 'June' if x==6 else 'July' if x==7 else 'Unknown')

plt.figure(figsize=(10,5))
bar_plot = sns.barplot(data=df, x='Month', y='Average_Delay_of_Month')
```

```
plt.xlabel('Month')
plt.ylabel('Average Delay per Month')
plt.show()
# insight: from data give, month with least average delay is January

{'_id': 1, 'avgDelay': 15.566352459016393}
{'_id': 2, 'avgDelay': 19.88471246959475}
{'_id': 3, 'avgDelay': 26.05433479824211}
```



26. Create a partitioning table “flights_partition” using partitioned by schema “Airports”

27. Write a MongoDB Map-Reduce function to calculate the total delay time for each airline, including all types of delays (e.g., AIR_SYSTEM_DELAY, SECURITY_DELAY, etc.).

28. Explain how you could use MongoDB’s Map-Reduce feature to find the most common origin-destination airport pairs with the longest delays.